

Forecasting Event Data

Zhanna Terechshenko

10/27/2017

Introduction

In this tutorial, I compare forecasting utility of PHOENIX and ICEWS event datasets using the output from random forest model.

Loading libraries:

```
rm(list=ls())
library(here)
library(caret)
library(stats)
library(plyr)
library(data.table)
library(tidyverse)
library(randomForest)
```

Making sure the results of this script are reproducible

```
set.seed(0927)
```

In order to compare forecasting utility I am using The Ground Truth Data Set (GTDS), which provides information on 5 events of interests (EOI): International Crisis, Ethnic/Religious Conflict, Domestic Crisis, Rebellion, and Insurgency. It covers 168 states from January 2001 through 2014. The unit is country-month.

```
eoi = read.csv("gtds_2001.to.may.2014.csv")
head(eoi)
```

```
##      ccode country year month   time ins reb dpc erv ic notes
## 1      20  Canada 2001     1 2001m1  0  0  0  0  0
## 2      20  Canada 2001     2 2001m2  0  0  0  0  0
## 3      20  Canada 2001     3 2001m3  0  0  0  0  0
## 4      20  Canada 2001     4 2001m4  0  0  0  0  0
## 5      20  Canada 2001     5 2001m5  0  0  0  0  0
## 6      20  Canada 2001     6 2001m6  0  0  0  0  0
##
##          coder insnotes dpcnotes rebnotes ervnotes icnotes
## 1 Bentley & Leonard
## 2 Bentley & Leonard
## 3 Bentley & Leonard
## 4 Bentley & Leonard
## 5 Bentley & Leonard
## 6 Bentley & Leonard
```

Here, I'll focus on International Crisis. For the purpose of forecasting I create an onset variable:

```
eoi_int = eoi %>%
  select(ccode, year, month, ic)

eoi_d = eoi_int[eoi_int$ic==1,]
eoi_d$onset <- with(eoi_d, ave(year, month, ccode, FUN = function(x)
  as.integer(c(TRUE, tail(x, -1L) != head(x, -1L) + 1L))))
eoi_d = eoi_d[eoi_d$onset==1,]
```

```
eoi_final = merge(eoi_int, eoi_d, all.x=T)
eoi_final$onset[is.na(eoi_final$onset)==T] <-0
eoi_final$ic = NULL
```

I've already preprocessed both PHOENIX and ICEWS. I aggregated both datasets to the country - month level and selected conflicts with government and military actors for international conflict. In this aggregated form, both datasets provide information on counts of events for quad classes: verbal cooperation, material cooperation, verbal conflict, material conflict.

```
phoenix_int = read.csv("pho_international.csv")
head(phoenix_int)
```

```
##   X ccode year month vcp mcp vcf mcf
## 1 1   100 2001     2   3   0   0   0
## 2 2   100 2001     4   1   0   0   0
## 3 3   100 2001     5   2   0   0   0
## 4 4   100 2001     6   6   0   0   0
## 5 5   100 2001     7   1   0   1   0
## 6 6   100 2001     9   1   0   0   1
```

I lagged quad categories by 3 and 6 months

```
pho_int_lag = phoenix_int %>%
  group_by(ccode) %>%
  arrange(ccode, year) %>%
  do(data.frame(., setNames(shift(.$vcp, c(3,6)), paste("vcp_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$mcp, c(3,6)), paste("mcp_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$vcf, c(3,6)), paste("vcf_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$mcf, c(3,6)), paste("mcf_1", c(3,6), sep=".")))) %>%
  select(ccode, year, month, vcp_1.3, mcp_1.3, vcf_1.3, mcf_1.3,
         vcp_1.6, mcp_1.6, vcf_1.6, mcf_1.6)

pho_int_lag = na.omit(pho_int_lag)
```

I merge GTDS and PHOENIX datasets together

```
df = merge(pho_int_lag, eoi_final, by=c("ccode", "year", "month"))
df$onset = as.factor(df$onset)
```

I split the data on training and test samples. In this example, I train the model on 2001 - 2005 data and test it on 2006

```
train_data = df[which(df$year<=2005),]
train_data_l3 = subset(train_data, select=c("ccode", "year", "month",
                                             "vcp_1.3", "mcp_1.3", "vcf_1.3", "mcf_1.3", 'onset'))
train_data_l6 = subset(train_data, select=c("ccode", "year", "month",
                                             "vcp_1.6", "mcp_1.6", "vcf_1.6", "mcf_1.6", 'onset'))

test_data = df[which(df$year==2006),]
test_data_l3 = subset(test_data, select=c("ccode", "year", "month",
                                             "vcp_1.3", "mcp_1.3", "vcf_1.3", "mcf_1.3", 'onset'))
test_data_l6 = subset(test_data, select=c("ccode", "year", "month",
                                             "vcp_1.6", "mcp_1.6", "vcf_1.6", "mcf_1.6", 'onset'))
```

I train the model using random forest and 5-fold cross validation.

```
rf_model<-train(onset~.,data=train_data_l3,method="rf",
               trControl=trainControl(method="cv",number=5, savePredictions = T),
               prox=TRUE,allowParallel=TRUE)
print(rf_model)
```

```
## Random Forest
##
## 5273 samples
##    7 predictor
##    2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 4219, 4219, 4217, 4218, 4219
## Resampling results across tuning parameters:
##
##  mtry  Accuracy  Kappa
##  2     0.9624513 0.3002987
##  4     0.9736396 0.6060027
##  7     0.9806565 0.7469173
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 7.
```

Looking at the results

```
print(rf_model$finalModel)
```


```
##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry, proximity = TRUE,      allowParallel = TRUE)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 7
##
##           OOB estimate of error rate: 1.54%
## Confusion matrix:
##      0   1 class.error
## 0 5020  26 0.005152596
## 1   55 172 0.242290749
```

```
testclass <- predict(rf_model, newdata = test_data_l3)
cfMatrix <- confusionMatrix(data = testclass, test_data_l3$onset)
print(cfMatrix)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    0    1
##           0 936  23
##           1  43   0
##
##           Accuracy : 0.9341
##           95% CI : (0.917, 0.9487)
##           No Information Rate : 0.977
##           P-Value [Acc > NIR] : 1.00000
```

```
##
##           Kappa : -0.0308
## Mcnemar's Test P-Value : 0.01935
##
##           Sensitivity : 0.9561
##           Specificity : 0.0000
##           Pos Pred Value : 0.9760
##           Neg Pred Value : 0.0000
##           Prevalence : 0.9770
##           Detection Rate : 0.9341
##           Detection Prevalence : 0.9571
##           Balanced Accuracy : 0.4780
##
##           'Positive' Class : 0
##
```

```
plot(rf_model)
```



```
test.probs <- predict(rf_model, test_data_l3, type="prob")

library(pROC)

pROC::roc(test_data_l3$onset, testclass)

rf.ROC <- roc(predictor=testclass,
               response=test_data_l3$onset,
               levels=rev(levels(test_data_l3$onset)))

rf.ROC$auc

plot(rf.ROC, main="RF ROC Phoenix")
```

I am doing the same for ICEWS data:

```
icw_int = read.csv("icw_international.csv")
head(icw_int)

##      X ccode year month  vcp mcp vcf mcf
## 1 1      2 2001      1 1377 11 27 13
## 2 2      2 2001      2 2189 10 69 14
## 3 3      2 2001      3 2431 23 58 32
## 4 4      2 2001      4 1710 7 34 16
## 5 5      2 2001      5 2385 9 52 13
## 6 6      2 2001      6 3032 27 71 18

# lag variables by 3 and 6 months
icw_int_lag = icw_int %>%
  group_by(ccode) %>%
  arrange(ccode, year) %>%
  do(data.frame(., setNames(shift(.$vcp, c(3,6)), paste("vcp_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$mcp, c(3,6)), paste("mcp_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$vcf, c(3,6)), paste("vcf_1", c(3,6), sep=".")))) %>%
  do(data.frame(., setNames(shift(.$mcf, c(3,6)), paste("mcf_1", c(3,6), sep=".")))) %>%
  select(ccode, year, month, vcp_1.3, mcp_1.3, vcf_1.3, mcf_1.3,
         vcp_1.6, mcp_1.6, vcf_1.6, mcf_1.6)

icw_int_lag = na.omit(icw_int_lag)

# merge 2 datasets
df2 = merge(icw_int_lag, eoi_final, by=c("ccode", "year", "month"))
df2$onset = as.factor(df2$onset)

# split the data
train_data = df2[which(df2$year<=2005),]
train_data_13 = subset(train_data, select=c("ccode", "year", "month",
      "vcp_1.3", "mcp_1.3", "vcf_1.3", "mcf_1.3", 'onset'))
train_data_16 = subset(train_data, select=c("ccode", "year", "month",
      "vcp_1.6", "mcp_1.6", "vcf_1.6", "mcf_1.6", 'onset'))

test_data = df2[which(df2$year==2006),]
test_data_13 = subset(test_data, select=c("ccode", "year", "month",
      "vcp_1.3", "mcp_1.3", "vcf_1.3", "mcf_1.3", 'onset'))
test_data_16 = subset(test_data, select=c("ccode", "year", "month",
      "vcp_1.6", "mcp_1.6", "vcf_1.6", "mcf_1.6", 'onset'))

# train the model
rf_model2<-train(onset~.,data=train_data_13,method="rf",
  trControl=trainControl(method="cv",number=5, savePredictions = T),
  prox=TRUE,allowParallel=TRUE)
print(rf_model2)

## Random Forest
##
```

```

## 7742 samples
## 7 predictor
## 2 classes: '0', '1'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 6194, 6194, 6194, 6193, 6193
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9711958 0.4508524
## 4 0.9814000 0.7083704
## 7 0.9855332 0.7843434
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 7.
print(rf_model2$finalModel)

##
## Call:
## randomForest(x = x, y = y, mtry = param$mtry, proximity = TRUE, allowParallel = TRUE)
## Type of random forest: classification
## Number of trees: 500
## No. of variables tried at each split: 7
##
## OOB estimate of error rate: 1.25%
## Confusion matrix:
## 0 1 class.error
## 0 7416 26 0.003493684
## 1 71 229 0.236666667


testclass <- predict(rf_model2, newdata = test_data_l3)
cfMatrix <- confusionMatrix(data = testclass, test_data_l3$onset)
print(cfMatrix)

## Confusion Matrix and Statistics
##
## Reference
## Prediction 0 1
## 0 1703 32
## 1 57 0
##
## Accuracy : 0.9503
## 95% CI : (0.9392, 0.9599)
## No Information Rate : 0.9821
## P-Value [Acc > NIR] : 1.00000
##
## Kappa : -0.0234
## McNemar's Test P-Value : 0.01096
##
## Sensitivity : 0.9676
## Specificity : 0.0000
## Pos Pred Value : 0.9816
## Neg Pred Value : 0.0000

```

```
##           Prevalence : 0.9821
##           Detection Rate : 0.9503
##           Detection Prevalence : 0.9682
##           Balanced Accuracy : 0.4838
##
##           'Positive' Class : 0
##
```

```
plot(rf_model2)
```



```
test.probs <- predict(rf_model2, test_data_l3, type="prob")
```

```
library(pROC)
```

```
pROC::roc(test_data_l3$onset, testclass)
```

```
rf.ROC <- roc(predictor=testclass,
               response=test_data_l3$onset,
               levels=rev(levels(test_data_l3$onset)))
```

```
rf.ROC$auc
```

```
# Area under the curve: 1
```

```
plot(rf.ROC, main="RF ROC ICEWS")
```